

# Applied Machine Learning!!!

W207 Section 9

Rasika Bhalerao

rasikabh@berkeley.edu

Aug 23: Welcome!  
Nov 8 and 22: No classes

# Schedule

## Supervised learning methods

	Sync	Topic
2	Aug 30	Linear Regression / Gradient Descent
3	Sep 6	Feature Engineering Bonus: Naive Bayes
4	Sep 13	Logistic Regression
5	Sep 20	Multiclass classification / Eval Metrics Bonus: Reinforcement learning
6	Sep 27	Neural Networks
7	Oct 4	KNN, Decision Trees, Ensembles

## Unsupervised learning methods

	Sync	Topic
8	Oct 11	KMeans and PCA Bonus: LDA
9	Oct 18	Text Embeddings Bonus: Language models
10	Oct 25	CNNs Bonus: GANs
11	Nov 1	EDA, Real data, Baselines
12	Nov 15	Fairness / Ethics
13	Nov 29	Fancy Neural Networks
14	Dec 6	Final Presentations

# Assignment Schedule

Due Date	Assignment
Aug 28	HW1
Sep 4	HW2
Sep 11	HW3
Sep 18	HW4
Sep 25	HW5
Oct 2	HW6
Oct 16	Group project baseline
Oct 23	HW8
Nov 6	HW9
Nov 20	HW10
Dec 4	Final project notebook + presentation

# Behavior expectations

- Healthy disagreement is expected
- Be mindful of one another's schedules
- Be a good listener
- Have fun in a professional manner
- Share related real-world experience
- Ask questions when something is confusing
- Keep it 100 but be respectful
- Be open-minded to new ideas in the real world and when coding
- On time for group meetings

Has anyone not signed up for a final  
project group?

[https://docs.google.com/document/d/1R3J\\_X1Rz6WP8eMQ2cyMC0wAr5iQdhMK\\_httdoNO6L0w/edit?usp=sharing](https://docs.google.com/document/d/1R3J_X1Rz6WP8eMQ2cyMC0wAr5iQdhMK_httdoNO6L0w/edit?usp=sharing)

# Softmax

What? Why?


$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

# Precision, Recall, and F1

	Ground truth positive	Ground truth negative
Tested positive	True positive (TP)	False positive (FP)
Tested negative	False negative (FN)	True negative (TN)

- Precision
  - Out of those tested positive, how many are truly positive?
  - $TP / (TP + FP)$
- Recall
  - Out of those truly positive, how many tested positive?
  - $TP / (TP + FN)$
- F1

$$\frac{2}{\text{recall}^{-1} + \text{precision}^{-1}}$$



How many selected items are relevant?

$$\text{Precision} = \frac{\text{green semi-circle}}{\text{green semi-circle} + \text{red semi-circle}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green semi-circle}}{\text{green semi-circle} + \text{green rectangle}}$$

# More than two classes

- **Multivalue classification:** each document can belong to 0, 1, or  $>1$  classes
- **Multinomial classification:** each document belongs to exactly 1 class
- Precision, recall, and F1 score are calculated for each class
  - **Macroaveraging:** get the scores for each class, then calculate the (unweighted) average
  - **Microaveraging:** count all the TP, FP, and FN for all classes and then calculate together
- “Accuracy” is simply the percent of documents classified correctly
  - Why is this less robust?



# Cross-entropy

A way to measure the difference between these two:

0.8	1
0.02	0
0.06	0
0.11	0
0.01	0

$$L = -\frac{1}{m} \sum_{i=1}^m y_i \cdot \log(\hat{y}_i)$$

# Confusion matrix

For each pair of classes  $\langle c_1, c_2 \rangle$ , how many documents from  $c_1$  were assigned to  $c_2$ ?

	Assigned positive	Assigned neutral	Assigned negative
True positive	8	3	2
True neutral	4	3	7
True negative	2	6	3

# Async Practice Quiz Questions (vote!)

As you adjust the classification threshold, when precision increases, recall increases too.	True	False
In a multiclass logistic regression model, each class has its own set of parameters.	True	False
Cross-entropy loss is only valid when the label distribution is one-hot.	True	False
Logistic regression output can be computed with a single matrix multiplication (and a sigmoid).	True	False
A linear model must rely on handcrafted input features if it is to understand complex structure.	True	False
A model trained for the MNIST digit classification task could be applied directly to mail sorting by zip code.	True	False

# Social bias in classification

No notebook this week (same as last week's binary classification code)

# Reinforcement learning

Y. LeCun

## How Much Information is the Machine Given during Learning?

### ► “Pure” Reinforcement Learning (**cherry**)

- The machine predicts a scalar reward given once in a while.


### ► **A few bits for some samples**

### ► Supervised Learning (**icing**)

- The machine predicts a category or a few numbers for each input
- Predicting human-supplied data
- **10→10,000 bits per sample**



### ► Self-Supervised Learning (**cake génoise**)

- The machine predicts any part of its input for any observed part.
- Predicts future frames in videos
- **Millions of bits per sample**



# Multi-Armed Bandit Problem


- Slot machine with  $n$  arms (bandits)
- Each arm has a different probability of giving you the reward
- Pulling an arm stochastically gives a reward (1) or failure (0)
- Task: maximize reward in the long run



<https://towardsdatascience.com/solving-the-multi-armed-bandit-problem-b72de40db97c>


# Which arm to pull?

- Pull all arms equally often?
- Only pull the arm that has given the best results so far?
- Mostly pull the “best” arm, but sometimes the others?
- An example of the exploration/exploitation dilemma



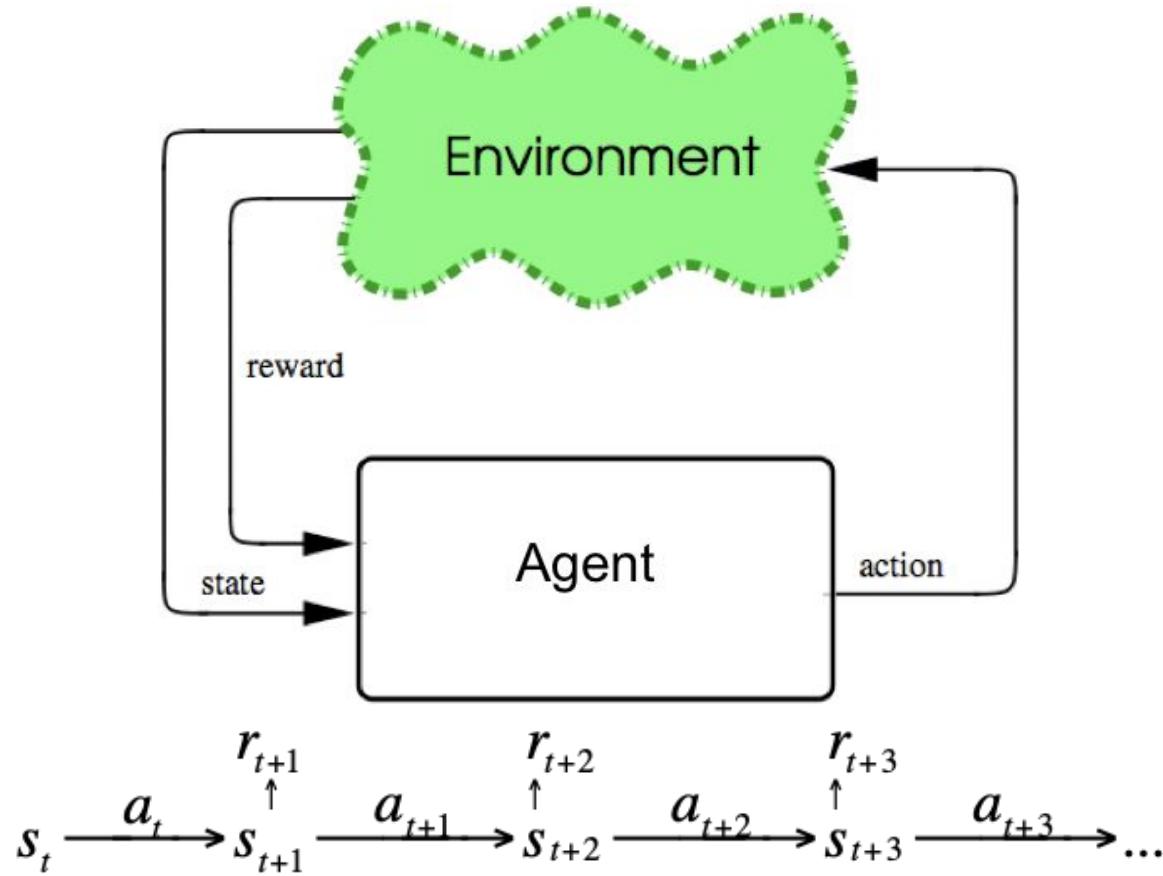
# What is reinforcement learning?

- Learning well-performing behavior from state observations and rewards
- Tree search:
  - Huge trees
  - State evaluation is hard
  - Each action selection may take a long time
- Supervised learning:
  - Agent is limited by expertise of expert from whom it learns
- Reinforcement learning:
  - Agent learns just from observations and rewards





# Reinforcement Learning



# Example rewards: PacMan

- One example:
  - 1 if you eat a pill
  - -10 if you get caught by a ghost
  - 2 if you eat a power pill or eat a ghost
  - 0 otherwise
- Another example:
  - -1 at every time step
  - 1,000,000 if you win the level



# Markov Decision Processes

- S: finite set of states (state space).  $s \in S$
- A: finite set of actions.  $a \in A$
- R: finite set of rewards.  $r \in R$
- State transition probabilities:

$$P_{ss'}^a = \Pr\{s_{t+1} = s' \mid s_t = s, a_t = a\}$$

- Policy:

$$\pi(s, a) = \Pr\{a_t = a \mid s_t = s\}$$

- Reward function:

$$R_{ss'}^a = E\{r_{t+1} \mid s_t = s, a_t = a\}$$

$s_t$  = state at time  $t$   
 $a_t$  = action chosen at time  $t$   
 $r_t$  = reward at time  $t$  (depends on time  $t-1$ )

# Markov Property

Probability of the next state and reward only depend on the immediately preceding state and action; it doesn't matter what happened before that

$$P_{ss'}^a = \Pr\{s_{t+1} = s', r_{t+1} = r \mid s_t, a_t, r_t, s_{t-1}, a_{t-1}, r_{t-1}, \dots, s_0, a_0, r_0\}$$

$$\equiv$$


$$P_{ss'}^a = \Pr\{s_{t+1} = s', r_{t+1} = r \mid s_t, a_t\}$$

# State Transition Probabilities

- Example with four states, considering one action
- Probability of going from one state to the other if the action is taken
- Each row sums to 1, outbound edges for each node sum to 1

	$s_1$	$s_2$	$s_3$	$s_4$
$s_1$	0	0.1	0.3	0.6
$s_2$	0.5	0	0	0.5
$s_3$	0	0.4	0	0.6
$s_4$	0	0.2	0.3	0.5

$$P_{ss'}^a$$



# Policy

- Determines agent's behavior
- Agent's goal: find the best policy
- Probability of taking action  $a$  when in state  $s$
- In general, the policy is stochastic:

$$\pi(s, a) = \Pr\{a_t = a \mid s_t = s\}$$

- If the policy is deterministic (we know which action to use) we can write:

$$\pi(s) \rightarrow a$$

# Reward Function

- Aka. expected return
- Agent's overall goal: find the policy that maximizes this
- Cumulative reward that the agent will receive from time  $t$  until the end of the game, if they take this action now

$$R_{ss'}^a = E\{r_{t+1} \mid s_t = s, a_t = a\}$$

- $\gamma$  close to 0  $\rightarrow$  agent cares more about immediate reward: shortsighted
- $\gamma$  close to 1  $\rightarrow$  agent cares more about future rewards: farsighted
- $\gamma = 1 \rightarrow$  sum doesn't converge for infinite time steps (fine if each episode always has finite steps)

$$R_t = \sum_{k=0}^T \gamma^k r_{t+k+1}, \quad 0 \leq \gamma \leq 1$$

# Value Function

Answers question: how “good” is this state or action?

Defined based on rewards expected from that state or action

Uses policy  $\pi$  to determine the value, given values of next states

- State-value:
  - What is the value of this state  $s$ , given policy  $\pi$ ?
- Action-value:
  - What is the value of taking action  $a$ , from state  $s$ , given policy  $\pi$ ?



# State-value function

Value of state  $s$  for an agent following policy  $\pi$

Expected return, starting state  $s$  at time  $t$  following policy  $\pi$

$$\begin{aligned} v_{\pi}(s) &= E_{\pi} [G_t \mid S_t = s] \\ &= E_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right] \end{aligned}$$

# Action-value function

Value of action  $a$  for an agent in state  $s$  following policy  $\pi$

Expected return, starting state  $s$  at time  $t$ , taking action  $a$ , then following policy  $\pi$

$$\begin{aligned} q_{\pi}(s, a) &= E_{\pi} [G_t \mid S_t = s, A_t = a] \\ &= E_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right] \end{aligned}$$

This is the Q-function, and its output is the Q-value!

# How does the agent find the best policy?

Optimal policy  $\pi$ : the policy with expected return  $\geq$  all other policies':

$$\pi \geq \pi' \text{ if and only if } v_{\pi}(s) \geq v_{\pi'}(s) \text{ for all } s \in \mathcal{S}$$

Based on that, the **optimal state-value function** for state  $s$  is:

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

And the **optimal action-value function** for state  $s$  is:

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

So,  $v^*$  is the largest expected return possible for each state  $s$

And  $q^*$  is the largest expected return possible for each state-action pair  $(s, a)$ .

# Bellman optimality equation for $q^*$

$q^*$  must satisfy the Bellman equation:

$$q_*(s, a) = E \left[ R_{t+1} + \gamma \max_{a'} q_*(s', a') \right]$$

Exercise: why is this true?

# Q-learning

- Objective: find the optimal policy by learning the optimal Q-values for each action-state pair
- Reminder: Q-function takes action-state pair, and returns expected return starting state  $s$  at time  $t$ , taking action  $a$ , then following policy  $\pi$

$$q_*(s, a) - q(s, a) = loss$$
$$E \left[ R_{t+1} + \gamma \max_{a'} q_*(s', a') \right] - E \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right] = loss$$

# Q-learning: updating Q-values

Learning rate  $\alpha$ : number between 0 and 1

Update Q-value  $q(s,a)$  with weighted sum of old and learned values

Higher learning rate  $\rightarrow$  more quickly adopt new Q-value

$$q^{new}(s, a) = (1 - \alpha) \underbrace{q(s, a)}_{\text{old value}} + \alpha \overbrace{\left( R_{t+1} + \gamma \max_{a'} q(s', a') \right)}^{\text{learned value}}$$

# Exploration vs. Exploitation

- Exploitation: take good actions in each state already taken before to maximize reward
- Exploration: take a chance on actions that may have lower value in order to learn more, and maybe find true best action to later exploit
- Need to balance the two!

# Exploration vs. Exploitation


- $\epsilon$ -greedy policy
  - Select greedy action  $1-\epsilon\%$  of the time (exploit), and a random action  $\epsilon\%$  of the time (explore)
  - $\epsilon$  decays over time
- Stochastic policy
  - Use action values to select actions probabilistically

$$\pi(s,b) = \frac{e^{Q(s,b)/\tau}}{\sum_a e^{Q(s,a)/\tau}}, \text{ where } \tau > 0 \text{ is the } \textit{temperature}$$

High temperatures increase exploration by making policy more random  
Low temperatures increase exploitation by making policy more greedy



# Q-learning example: The Lizard Game



Agent: lizard

Goal: Eat as many crickets as possible as fast as possible without meeting a bird


Actions: up, down, left, right

States: tiles

Rewards:

State	Reward	Game over?
1 cricket	1	No
Empty	-1	No
5 crickets	10	Yes
Bird	-10	Yes

# Q-learning example: The Lizard Game




Q-table:

States	Actions				
		Left	Right	Up	Down
	1 cricket	0	0	0	0
	Empty 1	0	0	0	0
	Empty 2	0	0	0	0
	Empty 3	0	0	0	0
	Bird	0	0	0	0
	Empty 4	0	0	0	0
	Empty 5	0	0	0	0
	Empty 6	0	0	0	0
5 crickets	0	0	0	0	

Update Q-values in this table

# Q-learning example: The Lizard Game



What would happen if we only did exploitation?

What would happen if we only did exploration?

State	Reward	Game over?
1 cricket	1	No
Empty	-1	No
5 crickets	10	Yes
Bird	-10	Yes

# Q-learning example: The Lizard Game


- First step: explore, move one tile to the right. Reward: -1
- Suppose discount rate  $\gamma = 0.99$  and learning rate  $\alpha = 0.7$

$$\begin{aligned} q^{new}(s, a) &= (1 - \alpha) \underbrace{q(s, a)}_{\text{old value}} + \alpha \overbrace{\left( R_{t+1} + \gamma \max_{a'} q(s', a') \right)}^{\text{new value}} \\ &= (1 - 0.7)(0) + 0.7 \left( -1 + 0.99 \left( \max_{a'} q(s', a') \right) \right) \end{aligned}$$

- To find maximum Q-value over all actions from  $s'$ , check table (currently all 0)

$$\begin{aligned} &= (1 - 0.7)(0) + 0.7(-1 + 0.99(0)) \\ &= 0 + 0.7(-1) \\ &= -0.7 \end{aligned}$$

# Q-learning example: The Lizard Game



States	Actions				
	Left	Right	Up	Down	
	1 cricket	0	0	0	0
	Empty 1	0	0	0	0
	Empty 2	0	0	0	0
	Empty 3	0	0	0	0
	Bird	0	0	0	0
	Empty 4	0	0	0	0
	Empty 5	0	-0.7	0	0
	Empty 6	0	0	0	0
5 crickets	0	0	0	0	